JOHNS HOPKINS
CAREY BUSINESS SCHOOL

Deep Learning and Natural Language Processing

Minghong Xu, PhD.
Associate Professor

Deep Learning Review

- Neural network models, inspired by neuroscience

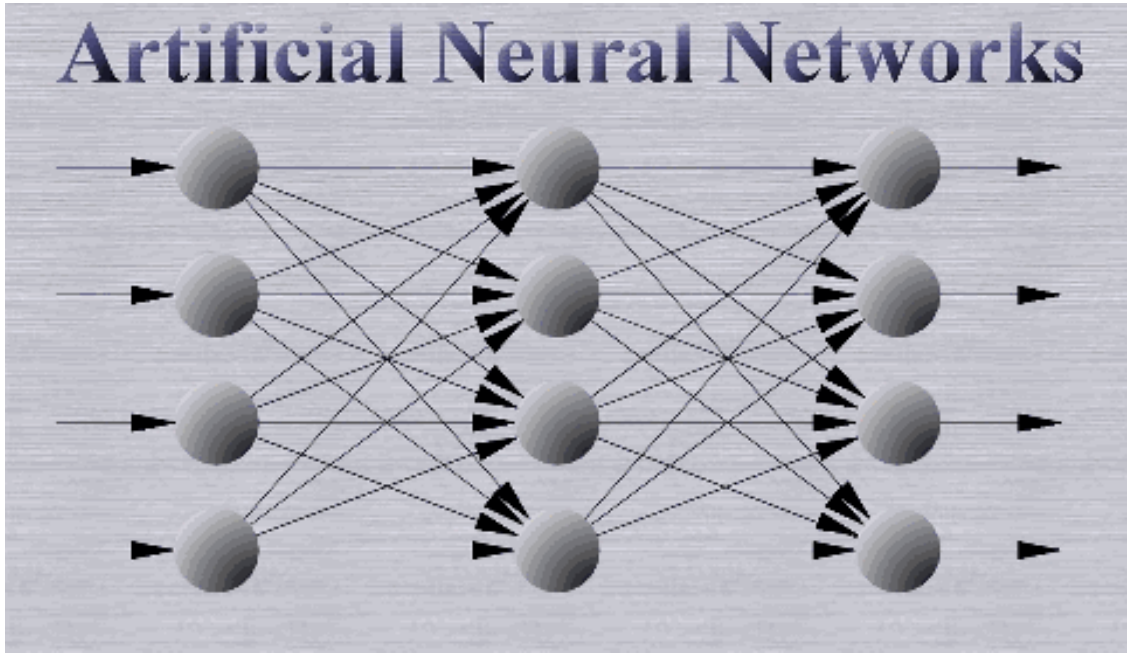

https://www.cse.unsw.edu.au/~cs9417ml/MLP2/

How does it work?

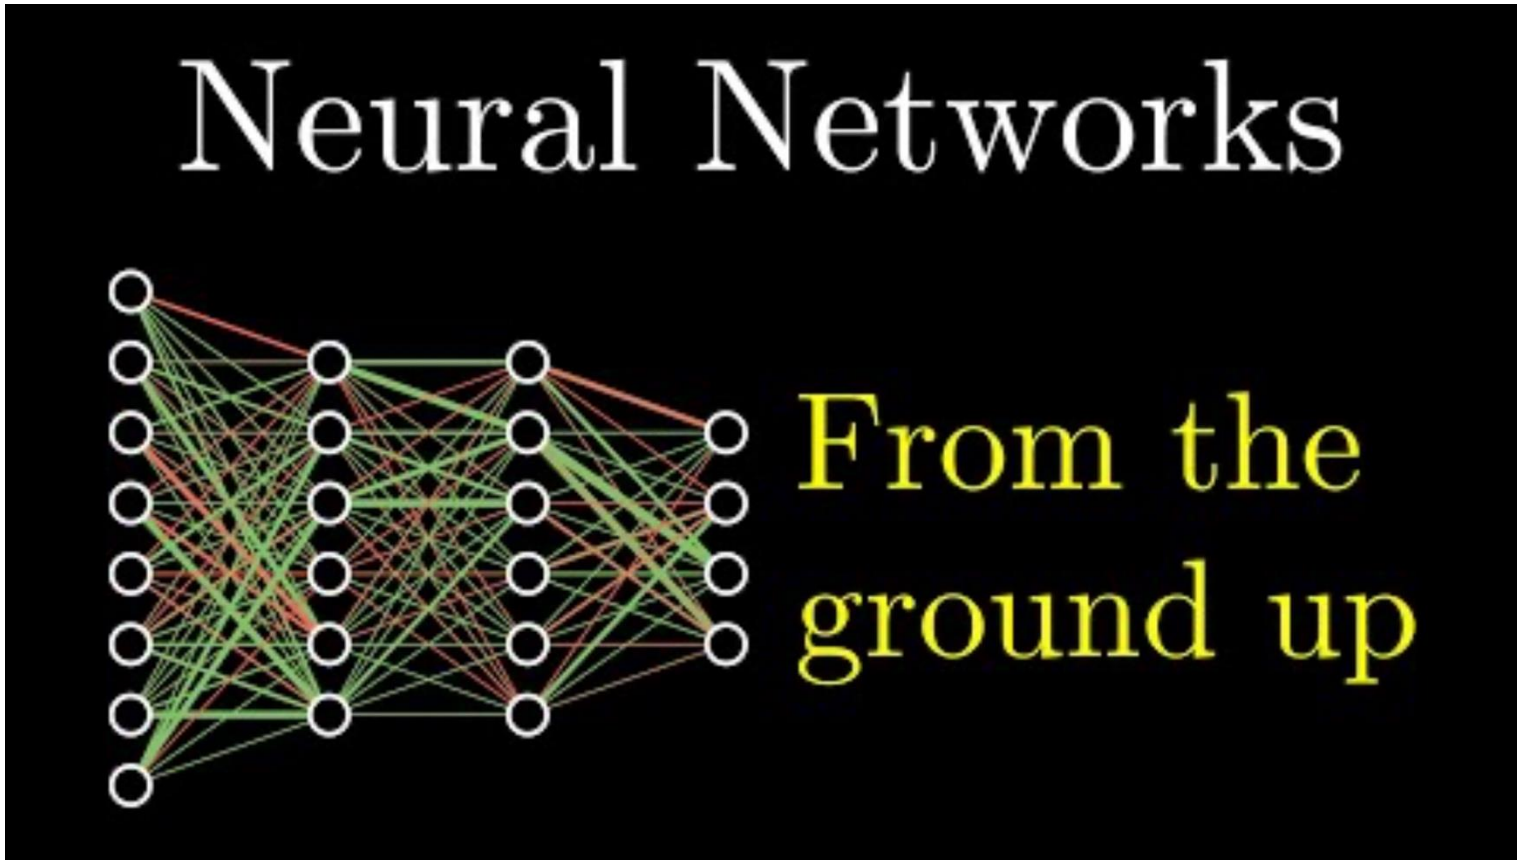

https://www.youtube.com/watch?v=aircAruvnKk&list=PLZHQObOWTQDNU6R1_67000Dx_ZCJB-3pi

Major Models: FFN

- Feedforward Neural Networks: plain models

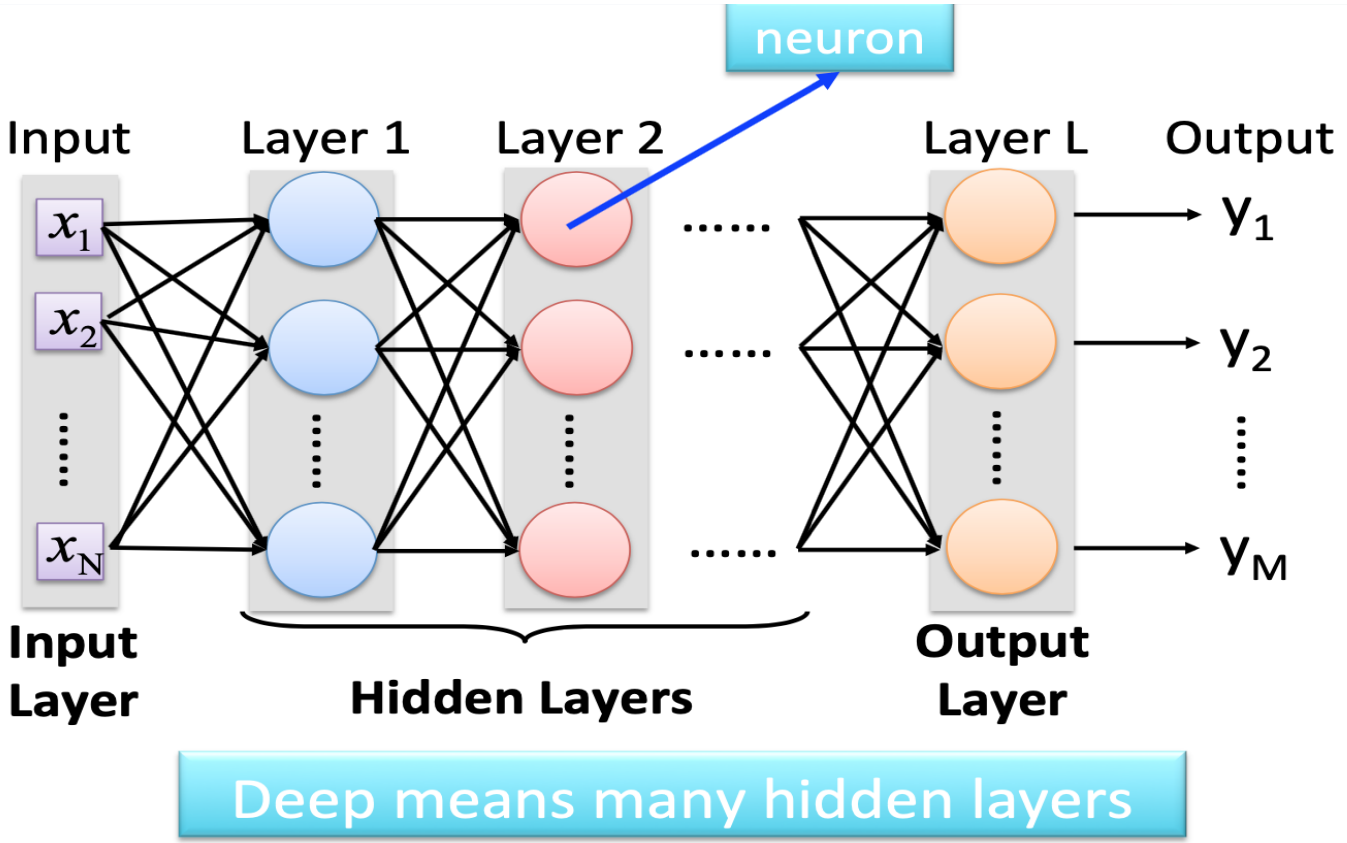

Major Models: CNN

- Convolutional Neural Networks: use convolution in some layers

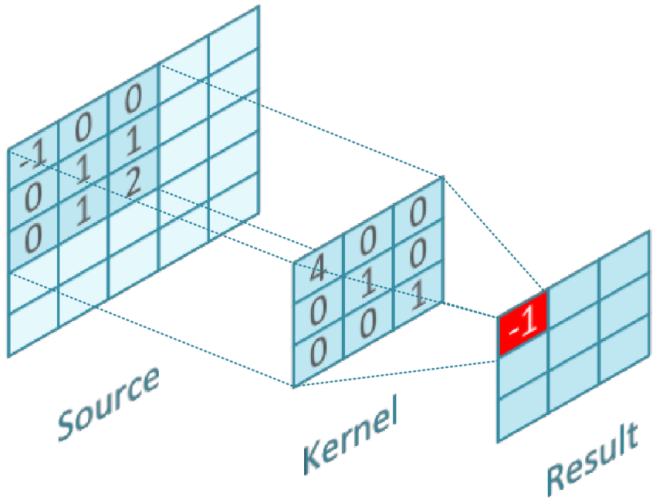

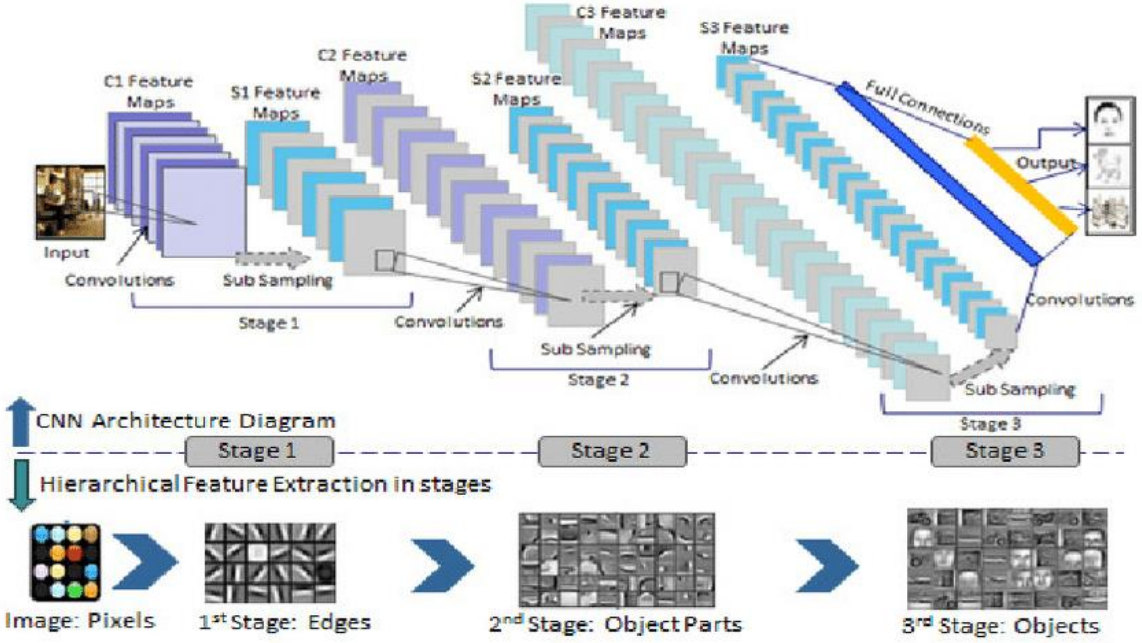

- https://setosa.io/ev/image-kernels/

Major Models: RNN

- Recurrent Neural Networks: the same network is used over time for sequential processing

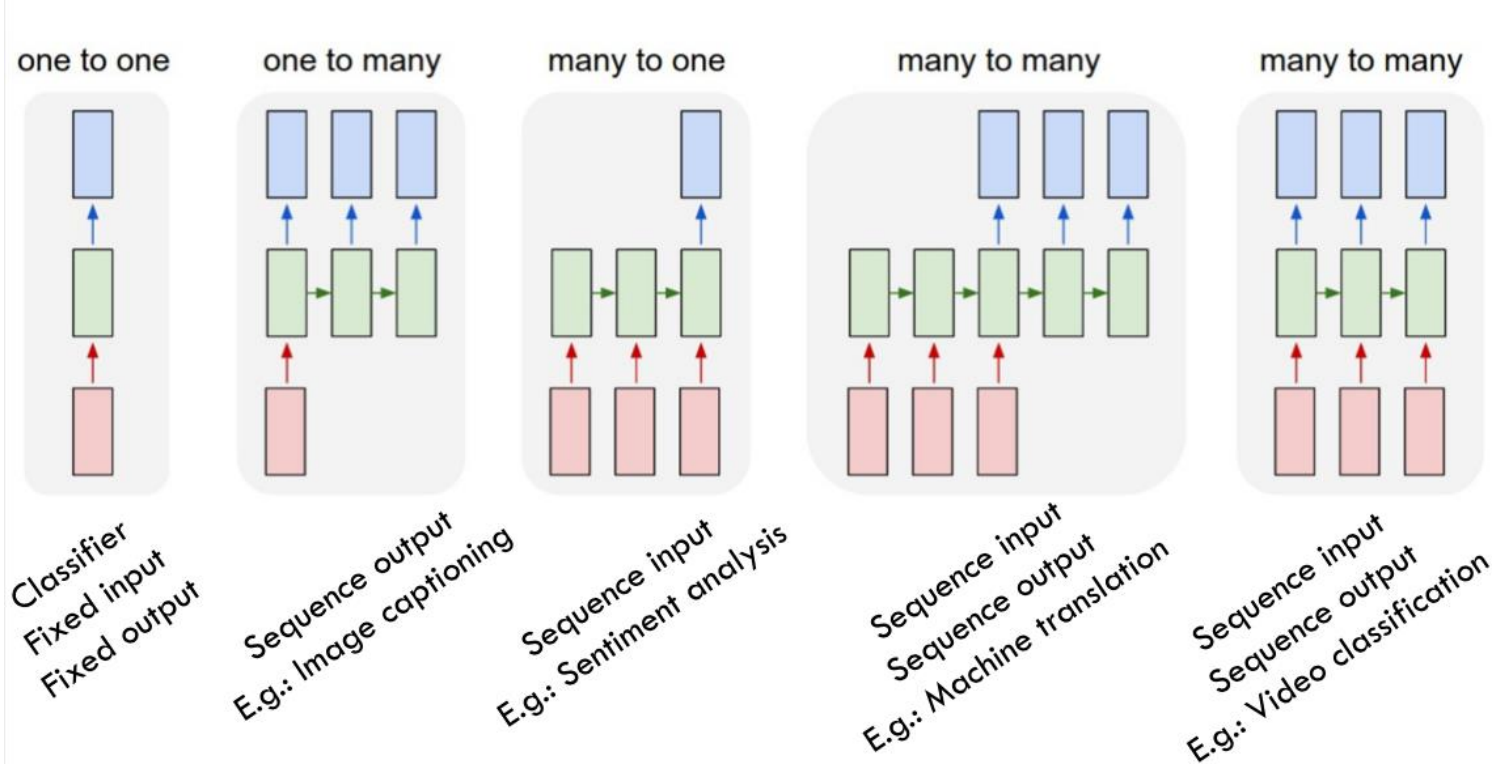

Natural Language Processing

What is NLP

- Text analytics, text mining
- Mine knowledge/information from huge amount of text data
- Many NLP systems are trained on very large collections of text (also called *corpora*)

Tokenization

- Break a stream of text into meaningful units
- Tokens: words, phrases, symbols

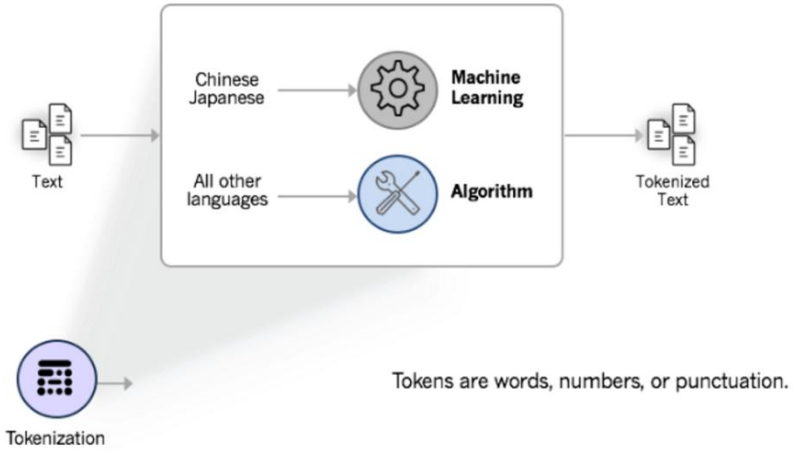

- https://platform.openai.com/tokenizer

Term Frequency

- Term: token in text
- A term is more important if it occurs more frequently in a document
- *So what to do?*
  - Count the occurrence!
- $tf(t,d) = frequency\ count\ of\ term\ t\ in\ doc\ d$
- This approach is called Bag of Words model

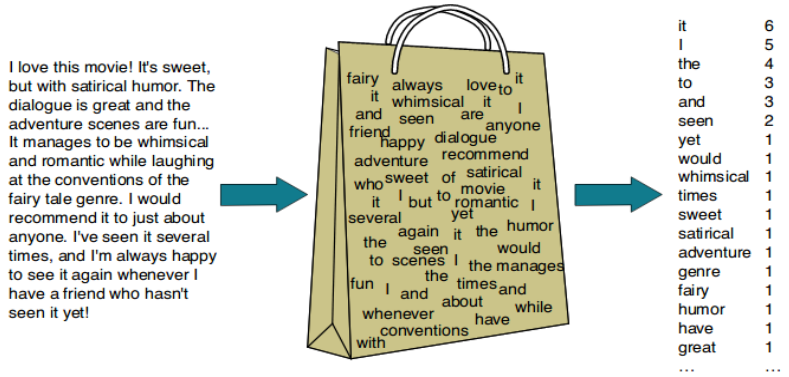

- *Any issue with this approach?*

TF Normalization

- Documents have different length
  - Doc 1 has 1000 words, and ‘coding’ appears 5 times
  - Doc 2 has 10 words, and ‘coding’ appears 2 times
- *How to solve this?*
  - Normalization!
- $tf(t,d) = \frac{frequency\ count\ of\ term\ t\ in\ doc\ d}{total\ words\ in\ doc\ d}$
- There are other ways to do normalization, such as maximum TF normalization

Stop Words

- ‘the’ ‘is’ are high frequency words in this document
- *But are they important?*
- Stop words: commonly used words in language
  - If it appears in all documents frequently, then it is not related with any specific doc
- *How can we identify them?*

Inverse Document Frequency

- Document frequency: a term is more discriminative if it occurs only in fewer document
- Inverse document frequency: assign higher weights to rare terms
- $idf(t) = \log(\frac{total\ documents}{documents\ with\ term\ t})$

- Combining *tf* and *idf*
- $tf \cdot idf = tf(t,d) \times idf(t)$

Lemmatization

- *How about “concepts” vs “concept”, “readings” vs “reading”?*
- *Another example can be “good” vs “best”*

- Lemmatization: the process of transforming a word into its root form

Vectorization

- Machine learning algorithms (including deep nets) require input to be vectors of numeric values
- Vectorization: represent words in a vector format

Word Embedding

- *Check the Token IDs on https://platform.openai.com/tokenizer*
- *Can we do better to represent each word?*
- Yes! We need the meaning

- Word Embedding:
  - Capture semantic meaning
  - Words that are semantically similar are close to each other in the vector space

Word2Vec

- Published in 2013 by researchers from Google

Efficient Estimation of Word Representations in Vector Space

Tomas Mikolov
Google Inc., Mountain View, CA
tmikolov@google.com

Kai Chen
Google Inc., Mountain View, CA
kaichen@google.com

Greg Corrado
Google Inc., Mountain View, CA
gcorrado@google.com

Jeffrey Dean
Google Inc., Mountain View, CA
jeff@google.com

Paper: Efficient Estimation of Word Representations in Vector Space

Two Examples

- “I will take a train from Baltimore to DC today.”

vs

- “I will train my first neural network model today.”

“You shall know a word by the company it keeps!”

by **John Rupert Firth**

Demo

- https://projector.tensorflow.org/
- Depend on training documents
- Statistical model
- Not symmetric

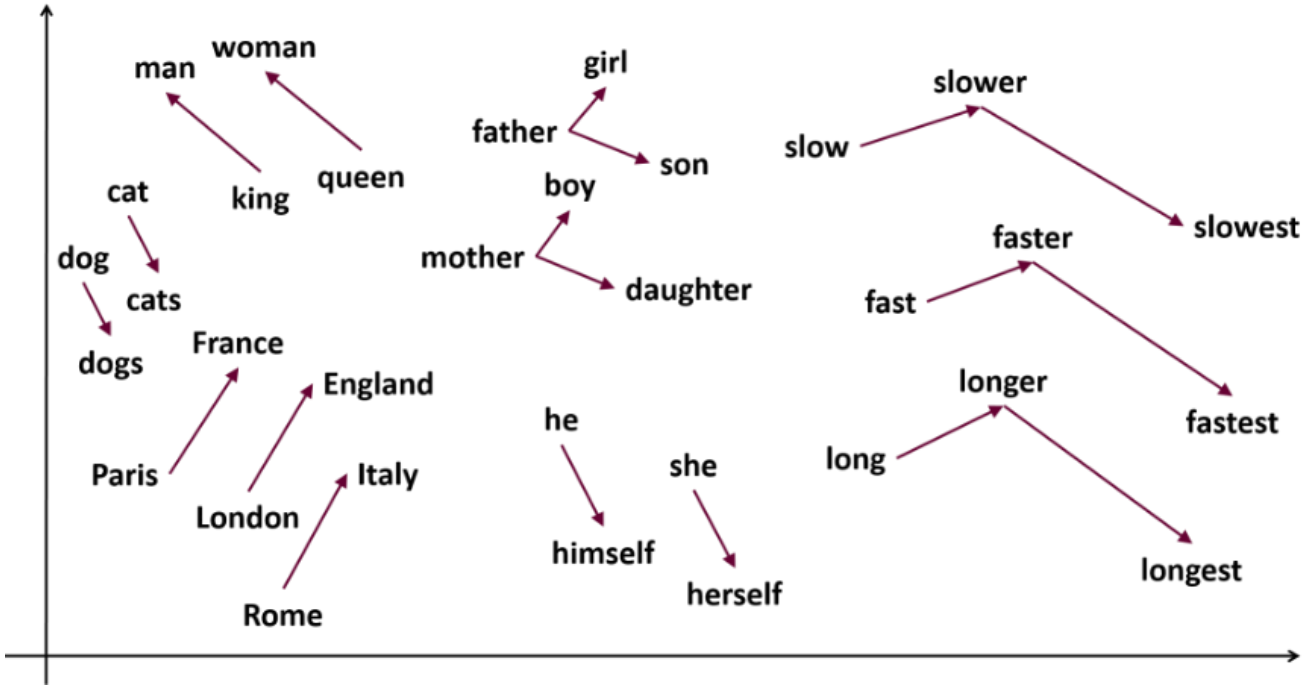

Colab Exercise on Word2Vec

- Download Word2vec.ipynb, ChatGPT_sentiment_txt.csv and ChatGPT_sentiment_txt_processed.csv from Canvas
- Upload the notebook to Colab
- Upload the two data files to your Google drive